

Problem_3(SVM)

April 21, 2020

```
[1]: # -*- coding: utf-8 -*-
      """
      Training an image classifier
      -----

      We will do the following steps in order:
      1. Load and normalizing the MNIST training and test datasets using
         ``torchvision``
      2. Define a SVM model
      3. Define a loss function
      4. Train the model on the training data
      5. Test the model on the test data
      1. Loading and normalizing MNIST
      ~~~~~~

      Using ``torchvision``, it's extremely easy to load MNIST.
      """

      import torch
      import torchvision
      import torchvision.transforms as transforms
      import itertools

      #####
      # The output of torchvision datasets are PILImage images of range [0, 1].
      # We transform them to Tensors of normalized range [-1, 1].
      # .. note::
      #     If running on Windows and you get a BrokenPipeError, try setting
      #     the num_worker of torch.utils.data.DataLoader() to 0.

[2]: pytorch_device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

[3]: transform = transforms.Compose(
      [transforms.ToTensor(),
       transforms.Normalize((0.1307,), (0.3081,))])

      trainset = torchvision.datasets.MNIST(root='./data', train=True,
                                             download=True, transform=transform)
      trainloader = torch.utils.data.DataLoader(trainset, batch_size=784,
                                                shuffle=True, num_workers=2)
```

```
testset = torchvision.datasets.MNIST(root='./data', train=False,
                                     download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=16,
                                          shuffle=False, num_workers=2)

classes = ('0', '1', '2', '3', '4', '5', '6', '7', '8', '9')
```

```
[4]: #####
# Let us show some of the training images, for fun.

import matplotlib.pyplot as plt
import numpy as np

# functions to show an image

def imshow(img):
    img = img / 2 + 0.5     # unnormalize
    npimg = img.cpu().numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)), cmap='gray')
    plt.show()

# get some random training images
examples = enumerate(trainloader)
batch_idx, (example_data, example_targets) = next(examples)

# show images
imshow(torchvision.utils.make_grid(example_data[:4]))

# print labels
print(' '.join('%5s' % classes[example_targets[j]] for j in range(4)))
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

<Figure size 640x480 with 1 Axes>

4 8 1 7

```
[6]: #####
# 2. Define a SVM model
# ~~~~~

import torch.nn as nn
import torch.nn.functional as F
```

```

class Model(nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.svm_layer = nn.Linear(784, 10) ### YOUR CODE ###, ### YOUR CODE ###

    def forward(self, x):
        # shape of input (=x): [16, 1, 28, 28]
        # shape of output: [16, 10]
        x = x.view(-1, 1 * 28 * 28)
        prediction = self.svm_layer(x) ### YOUR CODE ####
        return prediction

model = Model().to(pytorch_device)

```

```

[7]: #####
# 3. Define a Loss function and optimizer
# ~~~~~
# Let's use a Classification Cross-Entropy loss and SGD with momentum.

import torch.optim as optim

criterion = nn.CrossEntropyLoss().to(pytorch_device) ### YOUR CODE ###
optimizer = optim.SGD(model.parameters(), lr=0.01) ### YOUR CODE ###
#optimizer = optim.Adam(model.parameters(), lr=0.001)

```

```

[8]: #####
# 4. Train the model
# ~~~~~
#
# This is when things start to get interesting.
# We simply have to loop over our data iterator, and feed the inputs to the
# model and optimize.

for epoch in range(1000): # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs, labels = inputs.to(pytorch_device), labels.to(pytorch_device)

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = model(inputs)
        loss = criterion(outputs, labels) ### YOUR CODE ###

```

```

        loss.backward() ### YOUR CODE ###
        optimizer.step() ### YOUR CODE ###

        # print statistics
        running_loss += loss.item()
        if i % 2000 == 1999:    # print every 2000 mini-batches
            print('[%d, %5d] loss: %.3f' %
                  (epoch + 1, i + 1, running_loss / 2000))
            running_loss = 0.0

print('Finished Training')

```

Finished Training

```

[9]: #####
# See `here <https://pytorch.org/docs/stable/notes/serialization.html>`_
# for more details on saving PyTorch models.
#
# 5. Test the model on the test data
# ~~~~~
#
# We have trained the model for 2 passes over the training dataset.
# But we need to check if the model has learnt anything at all.
#
# We will check this by predicting the class label that the model
# outputs, and checking it against the ground-truth. If the prediction is
# correct, we add the sample to the list of correct predictions.
#
# Okay, first step. Let us display an image from the test set to get familiar.

dataiter = iter(testloader)
images, labels = dataiter.next()
images, labels = images.to(pytorch_device), labels.to(pytorch_device)

# print images
imshow(torchvision.utils.make_grid(images))

print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in
    ↪range(len(labels))))

#####
# Okay, now let us see what the model thinks these examples above are:

outputs = model(images)

#####
# The outputs are energies for the 10 classes.

```

```

# The higher the energy for a class, the more the model
# thinks that the image is of the particular class.
# So, let's get the index of the highest energy:
_, predicted = torch.max(outputs, 1)

print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
                                for j in range(len(labels))))

#####
# The results seem pretty good.
#
# Let us look at how the model performs on the whole dataset.

correct = 0
total = 0

with torch.no_grad():
    for data in trainloader:
        images, labels = data
        images, labels = images.to(pytorch_device), labels.to(pytorch_device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the model on the 60000 train images: %d %%' % (
    100 * correct / total))

correct = 0
total = 0

class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))
cmt = torch.zeros(10,10, dtype=torch.int64)

with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(pytorch_device), labels.to(pytorch_device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

        c = (predicted == labels).squeeze()
        for i in range(len(labels)):

```

```


        label = labels[i]
        cmt[labels[i], predicted[i]] += 1
        class_correct[label] += c[i].item()
        class_total[label] += 1

print('Accuracy of the model on the 10000 test images: %d %%' % (
    100 * correct / total))

for i in range(10):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))

```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



```

GroundTruth:      7      2      1      0      4      1      4      9      5      9      0
6      9      0      1      5
Predicted:        7      2      1      0      4      1      4      9      6      9      0
6      9      0      1      5
Accuracy of the model on the 60000 train images: 93 %
Accuracy of the model on the 10000 test images: 92 %
Accuracy of      0 : 97 %
Accuracy of      1 : 97 %
Accuracy of      2 : 90 %
Accuracy of      3 : 91 %
Accuracy of      4 : 93 %
Accuracy of      5 : 87 %
Accuracy of      6 : 95 %
Accuracy of      7 : 92 %
Accuracy of      8 : 88 %
Accuracy of      9 : 91 %

```

```

[10]: def plot_confusion_matrix(cm, classes, normalize=False, title='Confusion_
      ↪matrix', cmap=plt.cm.Blues):

```

```

if normalize:
    cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
    print("Normalized confusion matrix")
else:
    print('Confusion matrix, without normalization')

print(cm)
plt.imshow(cm, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt), horizontalalignment="center",
    ↪color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')

```

```

[11]: plt.figure(figsize=(10, 10))
      plot_confusion_matrix(cmt.numpy(), classes)

```

Confusion matrix, without normalization

```

[[ 960   0   1   3   0   7   4   4   1   0]
 [   0 1110   4   2   0   2   4   2  11   0]
 [   6   9 932  15   9   3  13   8  34   3]
 [   4   1  17 920   0  24   3  11  24   6]
 [   1   2   6   3 917   0   8   4   9  32]
 [  10   2   3  35  11 777  15   6  30   3]
 [   9   3   6   2   8  14 913   2   1   0]
 [   1   9  23   4   7   1   0 951   2  30]
 [   8  11   8  21   9  27  12   8 858  12]
 [  11   8   1   9  23   6   0  20   7 924]]

```

