

# Problem 4 (Final)

April 25, 2020

## 1 Problem 4

### 1.1 Reference

- <https://github.com/wihoho/Image-Recognition>
- <https://github.com/TrungTVo/spatial-pyramid-matching-scene-recognition>
- <https://stackoverflow.com/questions/52305578/sift-cv2-xfeatures2d-sift-create-not-working-even-though-have-contrib-instal>
- <https://github.com/CyrusChiu/Image-recognition> > Working environment
- Python 2.7
- Two versions of scikit-learn (0.20.1 and 0.18.1) was implemented. The detail reason is explained below its operational section.
- OpenCV (3.3.0.10) was implemented, because the 2nd reference recommends installing 2.xx version, but the 3rd reference said early 3.x version should be fine). The specific reason for this is the newer version of OpenCV does not provide SIFT anymore.
- Downloaded most of skeleton codes from the 4th reference, and its operational codes like as classifier.py, spm.py, utils.py from it, too.

```
[1]: import cv2
import spm
from spm import build_spatial_pyramid, spatial_pyramid_matching

import numpy as np
from sklearn import svm
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import SGDClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from mnist import MNIST

from utils import *
from classifier import svm_classifier
```

```
[2]: sklearn.__version__
```

```
[2]: '0.18.1'
```

```
[3]: all_data = MNIST('MNIST_data/')
```

```
[4]: train_input, train_label = all_data.load_training()
test_input, test_label = all_data.load_testing()
train_img_raw = np.zeros((len(train_input), len(train_input[0])))
test_img_raw = np.zeros((len(test_input), len(test_input[0])))

[5]: for i in range(len(train_input)):
    train_img_raw[i] = np.array(train_input[i]).astype('float32')
for i in range(len(test_input)):
    test_img_raw[i] = np.array(test_input[i]).astype('float32')

train_label = np.array(train_label).astype('uint8')
test_label = np.array(test_label).astype('uint8')

train_img_ratio = train_img_raw/255
test_img_ratio = test_img_raw/255

[6]: VOC_SIZE = 100
PYRAMID_LEVEL = 2

DSIFT_STEP_SIZE = 4
index_limit = 5000

# For the Test
train_img_raw = train_img_raw[:index_limit]
test_img_raw = test_img_raw[:5000]
train_label = train_label[:index_limit]
test_label = test_label[:5000]

x_train = [np.stack((train_img_raw[i].reshape((28,28)).
    ↪astype('uint8'), train_img_raw[i].reshape((28,28)).
    ↪astype('uint8'), train_img_raw[i].reshape((28,28)).astype('uint8')), axis=2)
    ↪for i in range(len(train_img_raw))]
y_train = train_label
x_test = [np.stack((test_img_raw[i].reshape((28,28)).
    ↪astype('uint8'), test_img_raw[i].reshape((28,28)).
    ↪astype('uint8'), test_img_raw[i].reshape((28,28)).astype('uint8')), axis=2)
    ↪for i in range(len(test_img_raw))]
y_test = test_label

[7]: # Dense SIFT feature extraction
x_train_feature = [extract_DenseSift_descriptors(img) for img in x_train]
x_test_feature = [extract_DenseSift_descriptors(img) for img in x_test]

# Remove None
train_tmp = [each for each in zip(x_train_feature, train_label) if not each[0]
    ↪is None]
x_train_feature, train_label = zip(*train_tmp)
```

```

test_tmp = [each for each in zip(x_test_feature, test_label) if not each[0] is_
↳None]
x_test_feature, test_label = zip(*test_tmp)

x_train_kp, x_train_des = zip(*x_train_feature)
x_test_kp, x_test_des = zip(*x_test_feature)

# print ("Train/Test split: {:d}/{:d}".format(len(y_train), len(y_test)))
# print ("Codebook Size: {:d}".format(VOC_SIZE))
# print ("Pyramid level: {:d}".format(PYRAMID_LEVEL))

# Building the codebook
#codebook = build_codebook(x_train_des, voc_size=VOC_SIZE)
"""
Because of the version confliction for scikit-learn, firstly we should make_
↳codebook file based on the latest version of
scikit-learn (0.20.1). The older version of this does not work. However, the_
↳latest version does not provide an
attribute, the grid_scores, which has mean_sores, scores, parameters as_
↳tuple-type. It is replaced with cv_results_, but
it gives lot of useless data here as dictionary type.
So, Once we process to make the codebook file with the latest version, then_
↳downgrade the scikit-learn to 0.18.1
, and proceed the following part.
"""
import cPickle
# with open('./bow_codebook.pkl', 'w') as f:
#     cPickle.dump(codebook, f)
with open('./bow_codebook.pkl', 'r') as f:
    codebook = cPickle.load(f)
    f.close()

```

Train/Test split: 5000/5000

Codebook Size: 100

Pyramid level: 2

```

[8]: # Spatial Pyramid Matching encoding
x_train = [spatial_pyramid_matching(x_train[i],
                                     x_train_des[i],
                                     codebook,
                                     level=PYRAMID_LEVEL)
            for i in range(len(x_train))]

x_test = [spatial_pyramid_matching(x_test[i],
                                    x_test_des[i],
                                    codebook,

```

```

                                level=PYRAMID_LEVEL) for i in
↪range(len(x_test))]

x_train = np.asarray(x_train)
x_test = np.asarray(x_test)

```

```

[9]: C_range = 10.0 ** np.arange(-3, 3)
gamma_range = 10.0 ** np.arange(-3, 3)
param_grid = dict(gamma=gamma_range.tolist(), C=C_range.tolist())

# Tuning hyperparameters
clf = GridSearchCV(svm.SVC(), param_grid, cv=3, n_jobs=-2)
clf.fit(x_train[:, y_train[:])

for params, mean_score, scores in clf.grid_scores_:
    print("%0.3f (+/-%0.03f) for %r" % (mean_score, scores.std() * 2, params))

y_true, y_pred = y_test, clf.predict(x_test)
print('\n')
# Printing report
print(classification_report(y_true, y_pred))

```

/Users/Haenara.SHIN/opt/anaconda3/envs/COGS260\_env\_27/lib/python2.7/site-packages/sklearn/model\_selection/\_search.py:667: DeprecationWarning: The grid\_scores\_ attribute was deprecated in version 0.18 in favor of the more elaborate cv\_results\_ attribute. The grid\_scores\_ attribute will not be available from 0.20

DeprecationWarning)

```

0.113 (+/-0.000) for {'C': 0.001, 'gamma': 0.001}
0.113 (+/-0.000) for {'C': 0.001, 'gamma': 0.01}
0.113 (+/-0.000) for {'C': 0.001, 'gamma': 0.1}
0.113 (+/-0.000) for {'C': 0.001, 'gamma': 1.0}
0.113 (+/-0.000) for {'C': 0.001, 'gamma': 10.0}
0.113 (+/-0.000) for {'C': 0.001, 'gamma': 100.0}
0.113 (+/-0.000) for {'C': 0.01, 'gamma': 0.001}
0.113 (+/-0.000) for {'C': 0.01, 'gamma': 0.01}
0.113 (+/-0.000) for {'C': 0.01, 'gamma': 0.1}
0.573 (+/-0.009) for {'C': 0.01, 'gamma': 1.0}
0.205 (+/-0.002) for {'C': 0.01, 'gamma': 10.0}
0.113 (+/-0.000) for {'C': 0.01, 'gamma': 100.0}
0.113 (+/-0.000) for {'C': 0.1, 'gamma': 0.001}
0.113 (+/-0.000) for {'C': 0.1, 'gamma': 0.01}
0.611 (+/-0.024) for {'C': 0.1, 'gamma': 0.1}
0.846 (+/-0.027) for {'C': 0.1, 'gamma': 1.0}
0.677 (+/-0.032) for {'C': 0.1, 'gamma': 10.0}
0.148 (+/-0.005) for {'C': 0.1, 'gamma': 100.0}
0.113 (+/-0.000) for {'C': 1.0, 'gamma': 0.001}

```

0.627 (+/-0.034) for {'C': 1.0, 'gamma': 0.01}  
 0.854 (+/-0.026) for {'C': 1.0, 'gamma': 0.1}  
 0.900 (+/-0.021) for {'C': 1.0, 'gamma': 1.0}  
 0.893 (+/-0.027) for {'C': 1.0, 'gamma': 10.0}  
 0.259 (+/-0.010) for {'C': 1.0, 'gamma': 100.0}  
 0.629 (+/-0.034) for {'C': 10.0, 'gamma': 0.001}  
 0.854 (+/-0.027) for {'C': 10.0, 'gamma': 0.01}  
 0.901 (+/-0.016) for {'C': 10.0, 'gamma': 0.1}  
 0.906 (+/-0.022) for {'C': 10.0, 'gamma': 1.0}  
 0.892 (+/-0.032) for {'C': 10.0, 'gamma': 10.0}  
 0.277 (+/-0.012) for {'C': 10.0, 'gamma': 100.0}  
 0.854 (+/-0.027) for {'C': 100.0, 'gamma': 0.001}  
 0.900 (+/-0.015) for {'C': 100.0, 'gamma': 0.01}  
 0.901 (+/-0.027) for {'C': 100.0, 'gamma': 0.1}  
 0.900 (+/-0.021) for {'C': 100.0, 'gamma': 1.0}  
 0.891 (+/-0.032) for {'C': 100.0, 'gamma': 10.0}  
 0.277 (+/-0.012) for {'C': 100.0, 'gamma': 100.0}

	precision	recall	f1-score	support
0	0.90	0.96	0.93	460
1	0.97	0.98	0.97	571
2	0.92	0.93	0.92	530
3	0.84	0.86	0.85	500
4	0.89	0.90	0.89	500
5	0.90	0.87	0.89	456
6	0.94	0.92	0.93	462
7	0.91	0.88	0.90	512
8	0.79	0.75	0.77	489
9	0.84	0.86	0.85	520
avg / total	0.89	0.89	0.89	5000

```
[10]: x_valid = x_train[:5000]
      y_valid = y_train[:5000]
      x_train = x_train[5000:]
      y_train = y_train[5000:]
```

```
[11]: valid_predictions=[]
      test_predictions=[]
      for i in range(1000):
          if i % 100 == 0:
              print('epoch ---> '+str(int(i/100)))
              valid_output = clf.predict([x_valid[i]])
              test_output = clf.predict([x_test[i]])
```

```

valid_predictions.append(valid_output)
test_predictions.append(test_output)

confusion_m = confusion_matrix(y_test[0:1000],test_predictions)

# Printing report
print (classification_report(y_test[0:1000],np.array(test_predictions)))

# Printing validation/test report (accuracy)
print ('validation accuracy is :',accuracy_score(y_valid[0:
↪1000],valid_predictions))
print ('test accuracy is :',accuracy_score(y_test[0:1000],test_predictions))

class_names = [chr(i) for i in range(ord('0'),ord('9')+1)]

```

```

epoch ---> 0
epoch ---> 1
epoch ---> 2
epoch ---> 3
epoch ---> 4
epoch ---> 5
epoch ---> 6
epoch ---> 7
epoch ---> 8
epoch ---> 9

```

	precision	recall	f1-score	support
0	0.94	0.98	0.96	85
1	0.99	0.98	0.99	126
2	0.89	0.94	0.92	116
3	0.93	0.86	0.89	107
4	0.92	0.89	0.91	110
5	0.94	0.92	0.93	87
6	0.94	0.97	0.95	87
7	0.91	0.94	0.93	99
8	0.80	0.74	0.77	89
9	0.83	0.90	0.87	94
avg / total	0.91	0.91	0.91	1000

```

('validation accuracy is :', 0.9869999999999999)
('test accuracy is :', 0.91400000000000003)

```

```

[12]: # Copied from previous problems
import matplotlib.pyplot as plt
import itertools

```

```

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

```

```

[13]: plot_confusion_matrix(confusion_m, classes=class_names, title='SPM confusion_
      ↪ matrix\invalid_acc='+str(accuracy_score(y_valid[0:1000], valid_predictions))+'\
      ↪ test_acc='+str(accuracy_score(y_test[0:1000], test_predictions)))

```

Confusion matrix, without normalization

```

[[ 83  0  0  0  1  0  0  0  1  0]
 [ 0 124  1  0  0  0  0  0  1  0]
 [ 0  0 109  2  0  0  1  3  1  0]
 [ 0  0  5 92  0  1  0  2  6  1]
 [ 0  1  1  0 98  0  2  0  0  8]
 [ 0  0  1  2  1 80  0  0  3  0]
 [ 2  0  0  0  1  0 84  0  0  0]
 [ 0  0  1  1  0  1  0 93  0  3]

```

```
[ 3  0  4  2  2  3  2  2 66  5]
[ 0  0  0  0  3  0  0  2  4 85]]
```

